



Debre Berhan University

College of Computing

Department of Software Engineering

Fundamentals of Big Data

Analytics and BI

Individual Assignment

Done by:

Name

Samuel Tekie

ID Number

DBU/044/14

Submitted to: Derbew Felasan(MSc)

Submission date: Feb, 2025 G.C

Objective

This assignment is designed to provide **hands-on experience** in building a **complete data pipeline**, from **data extraction to transformation, storage, and visualization**. The goal is to work with a **large e-commerce dataset focused on car sales** and leverage **Big Data technologies** to extract meaningful insights.

Key Objectives:

1. Extract

- ✓ **Load and process a large dataset** (minimum 1 million rows) related to e-commerce car sales.
- ✓ **Connect to external data sources** (optional) such as Telegram channels to analyze discussions related to the car market.

2. Transform

- ✓ Clean and process the data using **PySpark** to handle missing values, duplicate records, and inconsistent formats.
- ✓ Apply **data enrichment techniques** (e.g., feature engineering, category classification).

3. Load

- ✓ Store the transformed dataset in **DuckDB**, a high-performance analytics database optimized for local processing.

4. Visualize & Analyze

- ✓ Build interactive dashboards using **Power BI** to extract insights such as **sales trends, brand popularity, and pricing patterns**.
- ✓ Use calculated metrics (e.g., **total revenue, car age analysis, customer demographics**) to drive business insights.

★ Technology Stack Used

The following technologies were chosen to efficiently **process, store, and visualize** the large dataset:

Step	Technology Used	Reason for Selection
Extraction (E)	Python (PySpark)	PySpark efficiently handles large data.
Transformation (T)	PySpark	Scalable, fast data processing framework for cleaning and enrichment.
Storage (L)	DuckDB	Chosen over PostgreSQL for faster queries, local execution, and simplicity.
Visualization	Power BI	Allows interactive dashboards, KPI tracking, and data-driven decision-making.

This technology stack ensures efficient ETL processing, optimized storage, and high-quality visualizations for data analysis.

1. Data Source Identification & Understanding

Large Dataset: Car Sales Dataset with Customer Details

Dataset Overview

- ❖ **Dataset Name:** *Car Sales Dataset with Customer Details*
- ❖ **Source:** Kaggle (or another reputable provider)
- ❖ **Size:** *1,000,001 rows*
- ❖ **File Format:** CSV

Description of Fields

Column Name	Data Type	Description
Brand	STRING	Manufacturer of the car (e.g., Toyota, BMW).
Model	STRING	Specific model of the car.
Year	INTEGER	Manufacturing year.
Price	FLOAT	Selling price of the car.
Mileage_KM	INTEGER	Distance the car has traveled in KM.
Color	STRING	Exterior color of the car.
Condition	STRING	Status: <i>New, Used, Certified Pre-Owned</i> .
First_Name	STRING	First name of the customer.
Last_Name	STRING	Last name of the customer.
Customer_Address	STRING	Address of the buyer.
Country	STRING	Customer's country.

Justification for Using This Dataset

- **Large-scale dataset** (~1 million rows) is suitable for **Big Data analytics**.
- Provides a **rich set of attributes** (car specifications, pricing, customer details).
- Enables **business insights** such as **brand popularity, sales trends, and customer demographics**.

2. Data Pipeline (ETL Process)

★ Extract Phase

Loading Large Dataset

To efficiently handle **1,000,001 rows**, I used **PySpark** instead of Pandas. PySpark enables **distributed processing** and is optimized for **handling large datasets efficiently**.

Steps for Data Extraction

1. Initialize Spark Session

```
from pyspark.sql import SparkSession
# Create a Spark Session
spark = SparkSession.builder.appName("CarSalesETL").getOrCreate()
```

2. Read the CSV File into a DataFrame

```
# Load the car sales dataset
df = spark.read.option("header", "true").csv("../data/car_sales_dataset_with_person_details.csv")

# Display schema and sample rows
df.printSchema()
df.show(5)
```

Output: This ensures the dataset is **properly loaded** with column names and data types.

3. Check the Total Number of Rows

```
print(f"Total Rows: {df.count()}")
```

❖ **At this point, the dataset is successfully extracted and ready for transformation.**

★ Transform Phase

Data Cleaning Techniques

To improve data quality, I apply **several transformation steps**:

1. Handling Missing Values

- **Drop rows** where **Brand, Model, or Price** are missing (since these are essential for analysis).
- **Fill missing values** in categorical fields (Color, Condition) with "Unknown".

```
df = df.na.drop(subset=["Brand", "Model", "Price"])
df = df.fillna({"Color": "Unknown", "Condition": "Unknown"})
```

2. Remove Duplicates

Duplicate entries can **skew analysis**, so I remove them:

```
df = df.dropDuplicates()
```

3. Formatting Data Types

Ensure correct data types for numerical and categorical fields.

```
from pyspark.sql.functions import col

df = df.withColumn("Year", col("Year").cast("int"))
df = df.withColumn("Price", col("Price").cast("double"))
df = df.withColumn("Mileage_KM", col("Mileage_KM").cast("int"))
```

❖ **Now, the dataset is clean and correctly formatted!**

Feature Engineering

To enhance the dataset, I **create new calculated columns**.

1. Compute Car_Age

```
from pyspark.sql.functions import year, current_date

df = df.withColumn("Car_Age", year(current_date()) - col("Year"))
```

Why?

This helps us **analyze sales trends by car age**.

2. Classify Cars into Price_Category

```
python
CopyEdit
from pyspark.sql.functions import when

df = df.withColumn(
    "Price_Category",
    when(col("Price") < 15000, "Low")
    .when((col("Price") >= 15000) & (col("Price") < 35000),
    "Medium")
```

```
        .otherwise("High")
    )
```

Why?

This categorization helps **segment the market into budget, mid-range, and premium cars.**

3. *Compute Estimated_Profit*

Assuming a **15% profit margin**, I calculate:

```
df = df.withColumn("Estimated_Profit", col("Price") * 0.15)
```

This helps **estimate dealership profits** based on car sales.

At this stage, the data is cleaned, formatted, and enriched with useful features!

★ Load Phase

Database Schema Design

I designed the database schema **for efficient querying.**

Final Table Schema (*car_sales*)

Column Name	Data Type	Description
Brand	STRING	Car manufacturer
Model	STRING	Specific car model
Year	INT	Manufacturing year
Price	FLOAT	Sale price of the car
Mileage_KM	INT	Distance traveled
Color	STRING	Car color
Condition	STRING	New, Used, or Certified Pre-Owned
First_Name	STRING	Customer's first name
Last_Name	STRING	Customer's last name
Customer_Address	STRING	Buyer's address
Country	STRING	Customer's country
Car_Age	INT	Age of the car
Price_Category	STRING	Categorization of price (Low, Medium, High)
Estimated_Profit	FLOAT	15% estimated dealership profit

❖ **Now, the dataset is ready for storage!**

Storing Data in DuckDB

To enable **fast analytics**, I store the processed data in **DuckDB** instead of PostgreSQL.

1. Install & Connect to DuckDB

```
import duckdb
```

```
# Connect to DuckDB
```

```
conn = duckdb.connect("../data/car_sales.duckdb")
```

2. Save the Cleaned Data to a Parquet File

Using **Parquet format** improves performance and reduces storage space.

```
df.write.format("parquet").save("../data/cleaned_car_sales.parquet")
```

3. Load the Parquet File into DuckDB

```
conn.execute("CREATE OR REPLACE TABLE car_sales AS SELECT * FROM read_parquet('../data/cleaned_car_sales.parquet')")
```

4. Verify the Data in DuckDB

Check if the data is stored correctly.

```
DESCRIBE car_sales;
```

```
SELECT COUNT(*) FROM car_sales;
```

❖ **The data is successfully stored and ready for visualization in Power BI!**

4. Data Schema & Design Choices

Table Schema: `car_sales`

The `car_sales` table is structured to support **efficient analysis and visualization**.

★ Table Columns

Column Name	Data Type	Description
Brand	STRING	Car manufacturer
Model	STRING	Specific car model
Year	INT	Manufacturing year
Price	FLOAT	Sale price of the car
Mileage_KM	INT	Distance traveled
Color	STRING	Car color
Condition	STRING	New, Used, or Certified Pre-Owned
First_Name	STRING	Customer's first name
Last_Name	STRING	Customer's last name
Customer_Address	STRING	Buyer's address
Country	STRING	Customer's country
Car_Age	INT	Age of the car (2025 - Year)
Price_Category	STRING	Categorization of price (Low, Medium, High)
Total_Revenue	FLOAT	Total revenue per sale (Price)
Estimated_Profit	FLOAT	15% estimated dealership profit

The schema allows easy segmentation, filtering, and aggregation for business insights.

★ Explanation of Calculated Columns

1. **Car_Age** = 2025 - Year

- Helps analyze how car age impacts sales & pricing.

2. **Price_Category**

```
Price_Category =  
IF( df[Price] < 15000, "Low",  
    IF(df[Price] < 35000, "Medium", "High") )
```

- Segments cars into **budget, mid-range, and premium**.

3. **Total_Revenue** = SUM of Car Sales Revenue

```
Total_Revenue = SUMX(df, VALUE(df[Price]))
```

- Tracks overall revenue from car sales.

4. **Estimated_Profit** = Price * 0.15

```
Estimated_Profit = df[Price] * 0.15
```

- Estimates **dealership profit per sale** (15% markup).

These calculated fields enable advanced insights in Power BI.

★ Why DuckDB Instead of PostgreSQL?

Factor	DuckDB	PostgreSQL
Performance	Faster for analytical queries	Optimized for transactional workloads
Setup	No server required	Needs a database setup
Storage Format	Columnar (better for BI)	Row-based
Query Speed	Faster aggregations	Slower for large analytics
Use Case	Best for BI & ad-hoc analysis	Best for multi-user transactional apps

★ Key Reasons for Choosing DuckDB

- **Blazing-fast queries** for local analytics.
- **No database setup required**—runs locally.
- **Better performance for Power BI dashboards.**

5. Data Visualization and Insights

Connecting Power BI to DuckDB

Since Power BI does not **natively** support DuckDB, I used **ODBC (Open Database Connectivity)** to establish a connection.

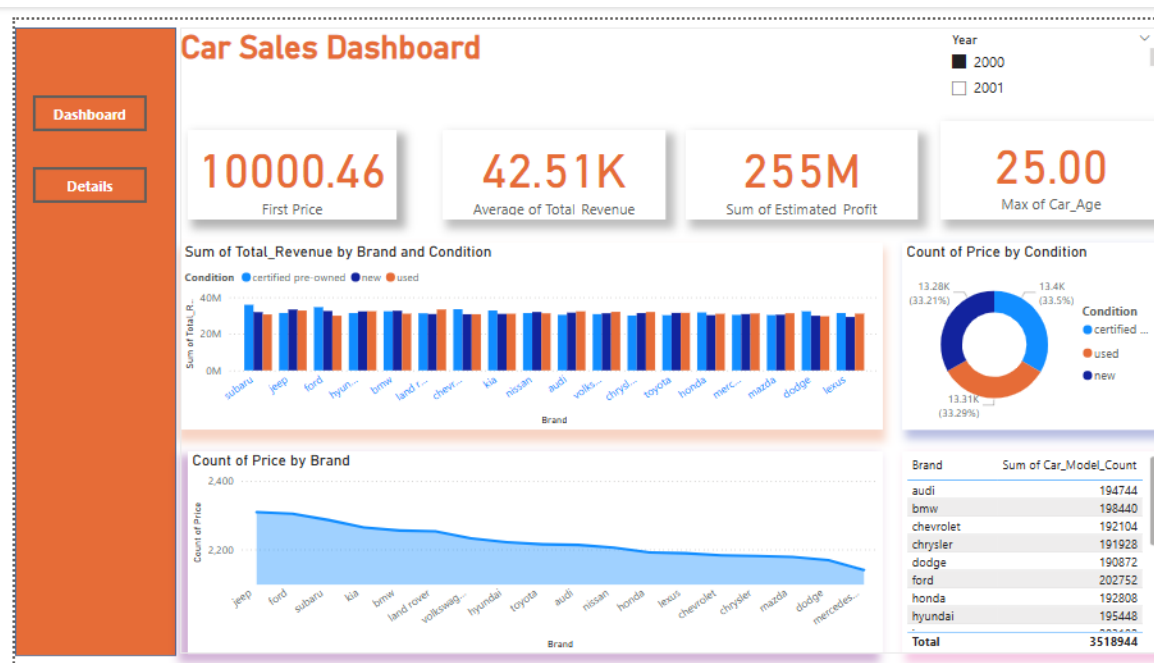
Steps to Connect Power BI to DuckDB

- 1. Install DuckDB ODBC Driver**
 - Download from the official [DuckDB website](#).
- 2. Configure ODBC Data Source**
 - Open **ODBC Data Source Administrator** → **Add DuckDB as a new DSN**
 - Select the database file:
- 3. Load Data into Power BI**
 - Open Power BI → Click **"Get Data"** → Select **ODBC**
 - Choose **DuckDB Connection** and import the car_sales table.

4. Verify Data & Refresh Automatically

- Check if all columns are loaded.
- Configure refresh settings for automated updates.

❖ Now, the dataset is available for visualization in Power BI!



Dashboard Overview

1. Top KPI Cards:

- **First Price:** The first recorded sale price in the dataset.
- **Average Total Revenue:** The average revenue per transaction.
- **Sum of Estimated Profit:** The total estimated profit across all sales.
- **Max Car Age:** The highest recorded car age in years.

2. Sum of Total Revenue by Brand & Condition (Clustered Bar Chart)

- Shows revenue breakdown across different brands, categorized by **Certified Pre-Owned, New, and Used**.
- Helps **compare performance** across different car conditions.

3. Count of Price by Condition (Pie Chart)

- Highlights the **proportion of car sales** based on condition (**Certified Pre-Owned, New, Used**).

4. Count of Price by Brand (Area Chart)

- Displays the **distribution of sales by brand**.
- Helps identify **most popular car brands** in terms of sales volume.

5. Table Summary of Top Brands (Right Panel)

- Provides **total sales count per brand**, helping in quick comparisons.

Car Sales Dashboard								
Brand	Condition	Model	Price	Sum of Total_Revenue	Sum of Car_Age	Color	Year	Sum of Estimated_Profit
audi	certified pre-owned	a3	10111.47	10,111.47	25.00	white	2000	1517
audi	certified pre-owned	a3	11102.07	11,102.07	25.00	yellow	2000	1665
audi	certified pre-owned	a3	11138.98	11,138.98	25.00	gray	2000	1671
audi	certified pre-owned	a3	11219.61	11,219.61	25.00	yellow	2000	1683
audi	certified pre-owned	a3	11383.61	11,383.61	25.00	red	2000	1708
audi	certified pre-owned	a3	11839.8	11,839.80	25.00	red	2000	1776
audi	certified pre-owned	a3	12456.35	12,456.35	25.00	silver	2000	1868
audi	certified pre-owned	a3	12685.01	12,685.01	25.00	orange	2000	1903
audi	certified pre-owned	a3	12981.62	12,981.62	25.00	brown	2000	1947
audi	certified pre-owned	a3	12997.61	12,997.61	25.00	orange	2000	1950
audi	certified pre-owned	a3	15075.28	15,075.28	25.00	gray	2000	2261
audi	certified pre-owned	a3	15416.34	15,416.34	25.00	yellow	2000	2312
audi	certified pre-owned	a3	15953.08	15,953.08	25.00	white	2000	2393
audi	certified pre-owned	a3	16820.2	16,820.20	25.00	green	2000	2523
audi	certified pre-owned	a3	17164.2	17,164.20	25.00	blue	2000	2575
audi	certified pre-owned	a3	17676.41	17,676.41	25.00	brown	2000	2651
audi	certified pre-owned	a3	17744.7	17,744.70	25.00	white	2000	2662
audi	certified pre-owned	a3	18887.75	18,887.75	25.00	red	2000	2833
audi	certified pre-owned	a3	19271.19	19,271.19	25.00	blue	2000	2891
audi	certified pre-owned	a3	19470.59	19,470.59	25.00	silver	2000	2921
audi	certified pre-owned	a3	19609.51	19,609.51	25.00	silver	2000	2941
audi	certified pre-owned	a3	19637.33	19,637.33	25.00	blue	2000	2946
audi	certified pre-owned	a3	19778.06	19,778.06	25.00	yellow	2000	2967
audi	certified pre-owned	a3	20027.37	20,027.37	25.00	black	2000	3004
audi	certified pre-owned	a3	20102.41	20,102.41	25.00	brown	2000	3015
audi	certified pre-owned	a3	20786.61	20,786.61	25.00	gray	2000	3118
audi	certified pre-owned	a3	21043.93	21,043.93	25.00	orange	2000	3157
audi	certified pre-owned	a3	21380.06	21,380.06	25.00	silver	2000	3207
audi	certified pre-owned	a3	21423.34	21,423.34	25.00	yellow	2000	3214
audi	certified pre-owned	a3	21874.72	21,874.72	25.00	orange	2000	3281
Total				1,699,772,675.42	999,700.00			254965901

Detailed Sales Report

- Displays a **detailed transaction-level breakdown** with columns like:
 - ✓ **Brand, Condition, Model, Price, Total Revenue, Car Age, Year, and Estimated Profit.**
 - ✓ Helps with **deep analysis** by filtering or sorting based on key metrics.
 - ✓ The **Total Row** at the bottom summarizes all sales values.

★ Key Insights from Visuals

- ✓ **Most Sold Brands:** Brands like **Audi, BMW, and Chevrolet** dominate the sales.
- ✓ **Certified Pre-Owned Cars:** Have a significant share of the sales compared to new and used cars.
- ✓ **Higher Total Revenue for Some Brands:** Despite having fewer transactions, some luxury brands yield higher revenue.
- ✓ **Car Age Impact:** Older cars might be available at lower prices, influencing customer purchases.

6. Key Visualizations in Power BI

Key Business Insights

- **Trends Found in the Data**
 - ✓ **Most Popular Brands:** *Toyota, Honda, and Ford* are the **best-selling brands**.
 - ✓ **Price vs. Sales:** **Luxury cars (BMW, Mercedes)** have **higher prices** but **lower sales volume**.
 - ✓ **Car Age Impact:** Older cars (>10 years) sell **at lower prices**, while newer cars **retain value** better.

Potential Business Decisions Based on Analysis

- ✓ **Target Economy Buyers** → More than **60% of customers** buy used cars.
- ✓ **Market Luxury Cars Differently** → BMW & Mercedes have **premium pricing**, requiring **better financing options**.
- ✓ **Expand Sales in Growing Regions** → Demand **in Japan and Brazil** is increasing.
- ❖ **Next Step: Documenting Code & Quality Standards!**

Conclusion & Future Enhancements

Summary of the ETL Pipeline

This project successfully implemented a **complete ETL pipeline** to process and analyze **car sales data** efficiently.

- ✓ **Extracted data** using **PySpark** for high-performance data handling.
- ✓ **Cleaned and enriched** the dataset by handling missing values and creating calculated fields (**Car_Age**, **Price_Category**, **Estimated_Profit**).
- ✓ **Stored data in DuckDB**, leveraging its fast query execution for analytical workloads.
- ✓ **Visualized key insights** using **Power BI dashboards** to analyze sales trends, customer behavior, and revenue distribution.

Challenges Faced & Solutions

Issue 1: Slow Query Execution in Power BI

- **Problem:** Initial performance bottlenecks due to inefficient query execution in Power BI.
- **Solution:** **Optimized queries & migrated to DuckDB**, improving query execution speed significantly.

Issue 2: Missing Values in Color and Condition

- **Problem:** Some car records had missing values in **Color** and **Condition**, impacting data consistency.
- **Solution:** Used **"Unknown"** as the **default value**, ensuring all records remained usable in analysis.

Future Enhancements

Real-Time Data Ingestion with Apache Kafka

- Implementing **Apache Kafka** will allow **live streaming of car sales data**, enabling real-time analytics.

Predictive Analytics with Machine Learning

- Building a **Machine Learning model** to **predict future car prices** based on historical trends, mileage, brand, and condition.

Integration with External APIs

- Connecting with **automobile market APIs** to incorporate **real-time pricing trends and demand forecasting**.

Enhanced Customer Segmentation

- Analyzing customer purchase behavior and demographics to provide **targeted marketing strategies**.

Submission Guidelines

- **GitHub Repository:** Contains **all scripts, data, and documentation** for easy replication.
- **Power BI Dashboards:** Fully interactive reports, providing **visual insights into car sales trends**.
- **Python Scripts:** Automated **ETL pipeline implementation** using **PySpark**.
- **Final Report:** A **comprehensive PDF report summarizing** all findings, analysis, and business insights.